



# KPLABS Course

HashiCorp Certified: Terraform Associate

## Domain 5 - Remote State Management

ISSUED BY

Zeal

REPRESENTATIVE

[instructors@kplabs.in](mailto:instructors@kplabs.in)

# Domain 5 - Remote State Management

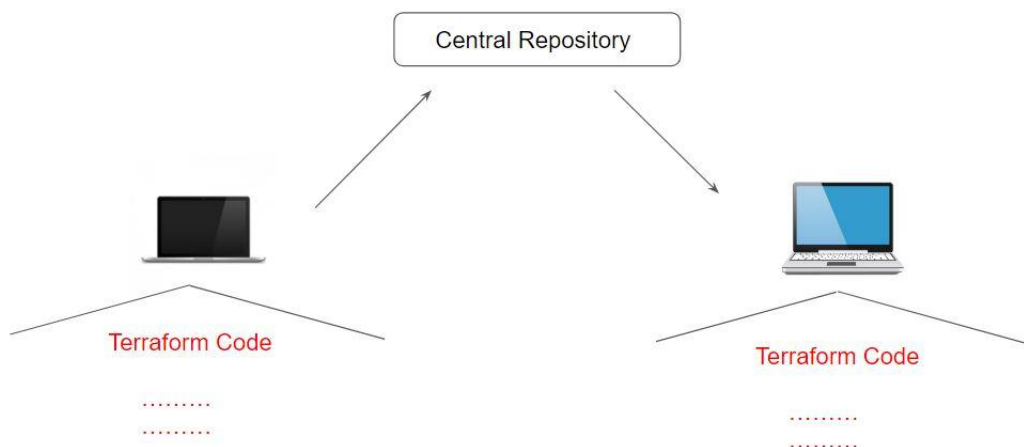
## Module 1: Integrating with GIT for team management

Till now, we have been working with terraform code locally.



However, storing your configuration files locally is not always an idea specifically in the scenario where other members of the team are also working on Terraform.

For such cases, it is important to store your Terraform code to a centralized repository like in Git.



## Module 2: Module Sources in Terraform

The source argument in a module block tells Terraform where to find the source code for the desired child module.

- Local paths
- Terraform Registry
- GitHub
- Bitbucket
- Generic Git, Mercurial repositories
- HTTP URLs
- S3 buckets
- GCS buckets

Let us explore some of the supported module sources.

### 2.1 Local Path

A local path must begin with either `./` or `../` to indicate that a local path is intended.

```
module "consul" {  
    source = "../consul"  
}
```

### 2.2 Git Module Source

Arbitrary Git repositories can be used by prefixing the address with the special `git::` prefix.

After this prefix, any valid Git URL can be specified to select one of the protocols supported by Git.

.

```

module "vpc" {
  source = "git::https://example.com/vpc.git"
}

module "storage" {
  source = "git::ssh://username@example.com/storage.git"
}

```

## 2.3 Referencing to a Branch

By default, Terraform will clone and use the default branch (referenced by HEAD) in the selected repository.

You can override this using the ref argument:

```

module "vpc" {
  source = "git::https://example.com/vpc.git?ref=v1.2.0"
}

```

The value of the ref argument can be any reference that would be accepted by the git checkout command, including branch and tag names.

## Module 3: Terraform & Gitignore

The .gitignore file is a text file that tells Git which files or folders to ignore in a project.

<b>.gitignore</b>
conf/
*.artifacts
credentials



Depending on the environment, it is recommended to avoid committing certain files to GIT.

Files to Ignore	Description
.terraform	This file will be recreated when terraform init is run.
terraform.tfvars	Likely to contain sensitive data like usernames/passwords and secrets.
terraform.tfstate	Should be stored in the remote side.
crash.log	If terraform crashes, the logs are stored to a file named crash.log

## Module 4: Terraform Backend

### 4.1 Basics of Backends

Backends primarily determine where Terraform stores its state.

By default, Terraform implicitly uses a backend called local to store state as a local file on disk.

```
provider "vault" {  
  address = "http://127.0.0.1:8200"  
}  
  
data "vault_generic_secret" "demo" {  
  path = "secret/db-creds"  
}  
  
output "vault_secrets" {  
  value = data.vault_generic_secret.demo.data_json  
  sensitive = "true"  
}
```

demo.tf

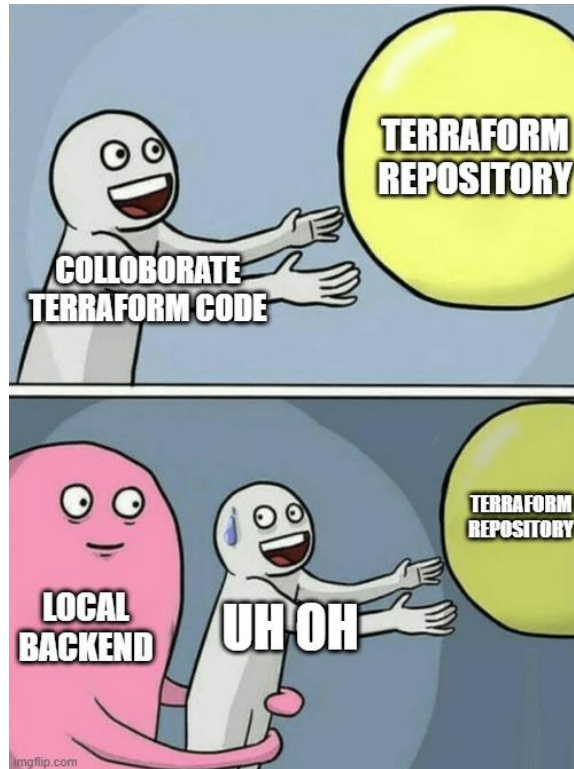
```
1 {  
2   "version": 4,  
3   "terraform_version": "1.1.9",  
4   "serial": 1,  
5   "lineage": "f7ba581a-ab47-b03e-2e54-e683a2dc4ba2",  
6   "outputs": {  
7     "vault_secrets": {  
8       "value": "{\"admin\":{\"password123\"}",  
9       "type": "string",  
10      "sensitive": true  
11    }  
12  },  
13  "resources": [  
14    {  
15      "mode": "data",  
16      "type": "vault_generic_secret",  
17      "name": "demo",  
18      "provider": "provider[\"registry.terraform.io/hashicorp/vault\"]",  
19      "instances": [  
20
```

terraform.tfstate

### 4.2 Challenge with Local Backend

Nowadays Terraform projects are handled and collaborated by an entire team.

Storing the state file in the local laptop will not allow collaboration.




### 4.3 Ideal Architecture

Following describes one of the recommended architectures:

The Terraform Code is stored in Git Repository.

The State file is stored in a Central backend.



#### 4.4 Backends Supported in Terraform

Terraform supports multiple backends that allow remote service related operations.


Some of the popular backends include:

- S3
- Consul
- Azurearm
- Kubernetes
- HTTP
- ETCD

#### 4.5 Important Note

Accessing state in a remote service generally requires some kind of access credentials

Some backends act like plain "remote disks" for state files; others support locking the state while operations are being performed, which helps prevent conflicts and inconsistencies.



## Module 5: State File Locking

### 5.1 Understanding State Lock

Whenever you are performing write operation, terraform would lock the state file.

This is very important as otherwise during your ongoing terraform apply operations, if others also try for the same, it can corrupt your state file.

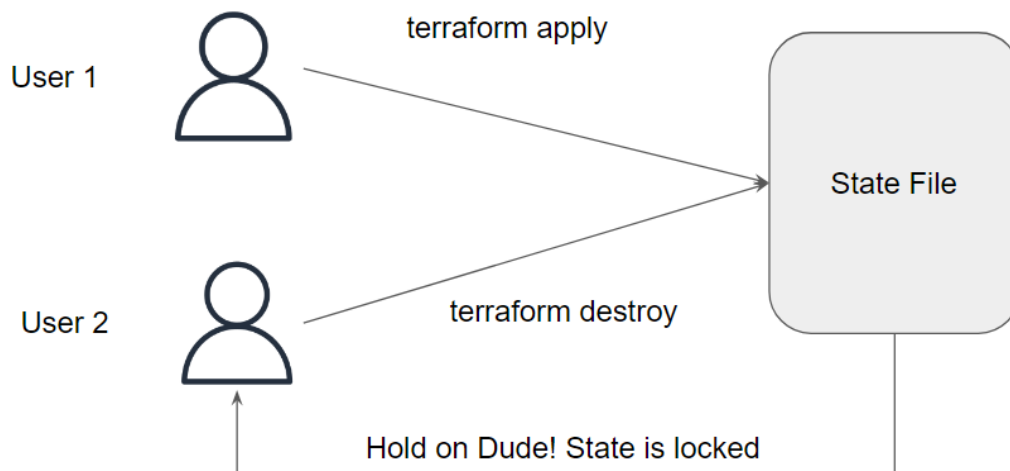
```
C:\Users\Zeal Vora\Desktop\tf-demo\remote-backend>terraform plan
```

```
Error: Error acquiring the state lock
```

```
Error message: Failed to read state file: The state file could not be read: read terraform.tfstate: The process cannot access the file because another process has locked a portion of the file.
```

```
Terraform acquires a state lock to protect the state from being written by multiple users at the same time. Please resolve the issue above and try again. For most commands, you can disable locking with the "-lock=false" flag, but this is not recommended.
```

Following diagram illustrates the basic working.





## 5.2 Important Note

State locking happens automatically on all operations that could write state. You won't see any message that it is happening

If state locking fails, Terraform will not continue

Not all backends support locking. The documentation for each backend includes details on whether it supports locking or not.

## 5.3 Force Unlocking State

Terraform has a **force-unlock** command to manually unlock the state if unlocking failed.


If you unlock the state when someone else is holding the lock it could cause multiple writers.

Force unlock should only be used to unlock your own lock in the situation where automatic unlocking failed.

## Module 6: State Locking in S3 Backend

By default, S3 does not support State Locking functionality.

You need to make use of the DynamoDB table to achieve state locking functionality.



## Module 7: Terraform State Management

As your Terraform usage becomes more advanced, there are some cases where you may need to modify the Terraform state.

It is important to never modify the state file directly. Instead, make use of terraform state command.

There are multiple sub-commands that can be used with terraform state, these include:

State Sub Command	Description
list	List resources within terraform state file.
mv	Moves item with terraform state.
pull	Manually download and output the state from remote state.
push	Manually upload a local state file to remote state.
rm	Remove items from the Terraform state
show	Show the attributes of a single resource in the state.

### 7.1 Sub Command - List

The terraform state list command is used to list resources within a Terraform state.

```
bash-4.2# terraform state list
aws_iam_user.lb
aws_instance.webapp
```

### 7.2 Sub Command - Move

The terraform state mv command is used to move items in a Terraform state.

This command is used in many cases in which you want to rename an existing resource without destroying and recreating it.

Due to the destructive nature of this command, this command will output a backup copy of the state prior to saving any changes

Overall Syntax:

```
terraform state mv [options] SOURCE DESTINATION
```

### 7.3 Sub Command - Pull

The terraform state pull command is used to manually download and output the state from a remote state.

This is useful for reading values out of state (potentially pairing this command with something like jq).

### 7.4 Sub Command - Push

The terraform state push command is used to manually upload a local state file to remote state.

This command should rarely be used.

### 7.5 Sub Command - Remove

The terraform state rm command is used to remove items from the Terraform state.

Items removed from the Terraform state are not physically destroyed.

Items removed from the Terraform state are only no longer managed by Terraform

For example, if you remove an AWS instance from the state, the AWS instance will continue running, but terraform plan will no longer see that instance.

## 7.6 Sub Command - Show

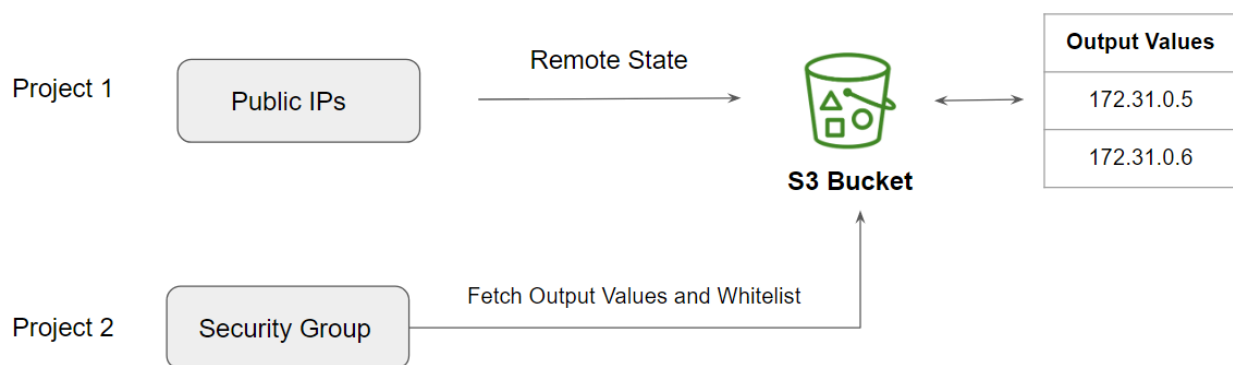
The terraform state show command is used to show the attributes of a single resource in the Terraform state.

```
bash-4.2# terraform state show aws_instance.webapp
# aws_instance.webapp:
resource "aws_instance" "webapp" {
  ami                = "ami-082b5a644766e0e6f"
  arn                = "arn:aws:ec2:us-west-2:018721151861:instance/i-0107ea9ed06c467e0"
  associate_public_ip_address = true
  availability_zone   = "us-west-2b"
  cpu_core_count      = 1
  cpu_threads_per_core = 1
  disable_api_termination = false
  ebs_optimized        = false
  get_password_data    = false
  id                  = "i-0107ea9ed06c467e0"
  instance_state      = "running"
  instance_type       = "t2.micro"
```

## Module 8: Connecting Remote States

### 8.1 Basics of Terraform Remote State

The [terraform\\_remote\\_state](#) data source retrieves the root module output values from some other Terraform configuration, using the latest state snapshot from the remote backend.



## 8.2 Step 1 - Create a Project with Output Values & S3 Backend

```
resource "aws_eip" "lb" {
  vpc      = true
}

output "eip_addr" {
  value = aws_eip.lb.public_ip
}
```



```
terraform {
  backend "s3" {
    bucket = "kplabs-terraform-backend"
    key    = "network/eip.tfstate"
    region = "us-east-1"
  }
}
```

## 8.3 Step 2 - Reference Output Values from Different Project

```
data "terraform_remote_state" "eip" {
  backend = "s3"
  config = {
    bucket = "kplabs-terraform-backend"
    key    = "network/eip.tfstate"
    region = "us-east-1"
  }
}
```



```
resource "aws_security_group" "allow_tls" {
  name        = "allow_tls"
  description = "Allow TLS inbound traffic"

  ingress {
    description = "TLS from VPC"
    from_port   = 443
    to_port     = 443
    protocol    = "tcp"
    cidr_blocks = ["${data.terraform_remote_state.eip.outputs.eip_addr}/32"]
  }
}
```

## Module 9: Terraform Import

It might happen that there is a resource that is already created manually.

In such a case, any change you want to make to that resource must be done manually.



web.tf

Terraform is able to import existing infrastructure. This allows you to take resources you've created by some other means and bring it under Terraform management.

The current implementation of Terraform import can only import resources into the state. It does not generate configuration. A future version of Terraform will also generate configuration.

Because of this, prior to running terraform import it is necessary to write manually a resource configuration block for the resource, to which the imported object will be mapped.

## Join Our Discord Community

We invite you to join our Discord community, where you can interact with our support team for any course-based technical queries and connect with other students who are doing the same course.

### Joining URL:

<http://kplabs.in/chat>

